

Classification (1)

POSTECH MIP Lab.

TA: Joonhyuk Park, Seunghun Baek, Soojin Hwang

Goal

- Understand classification task
- Understand the development of CNN based classification models (**AlexNet**, **VGG**, GoogLeNet, ResNet)
- Learn how to implement CNN models from architecture

Classification task

Classification



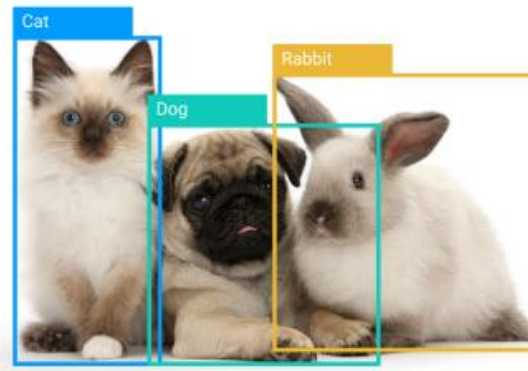
Cat

Classification
+Localization




Cat

Object
Detection



Cat, Dog, Rabbit

Instance
Segmentation



Cat, Dog, Rabbit

Classification task




```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

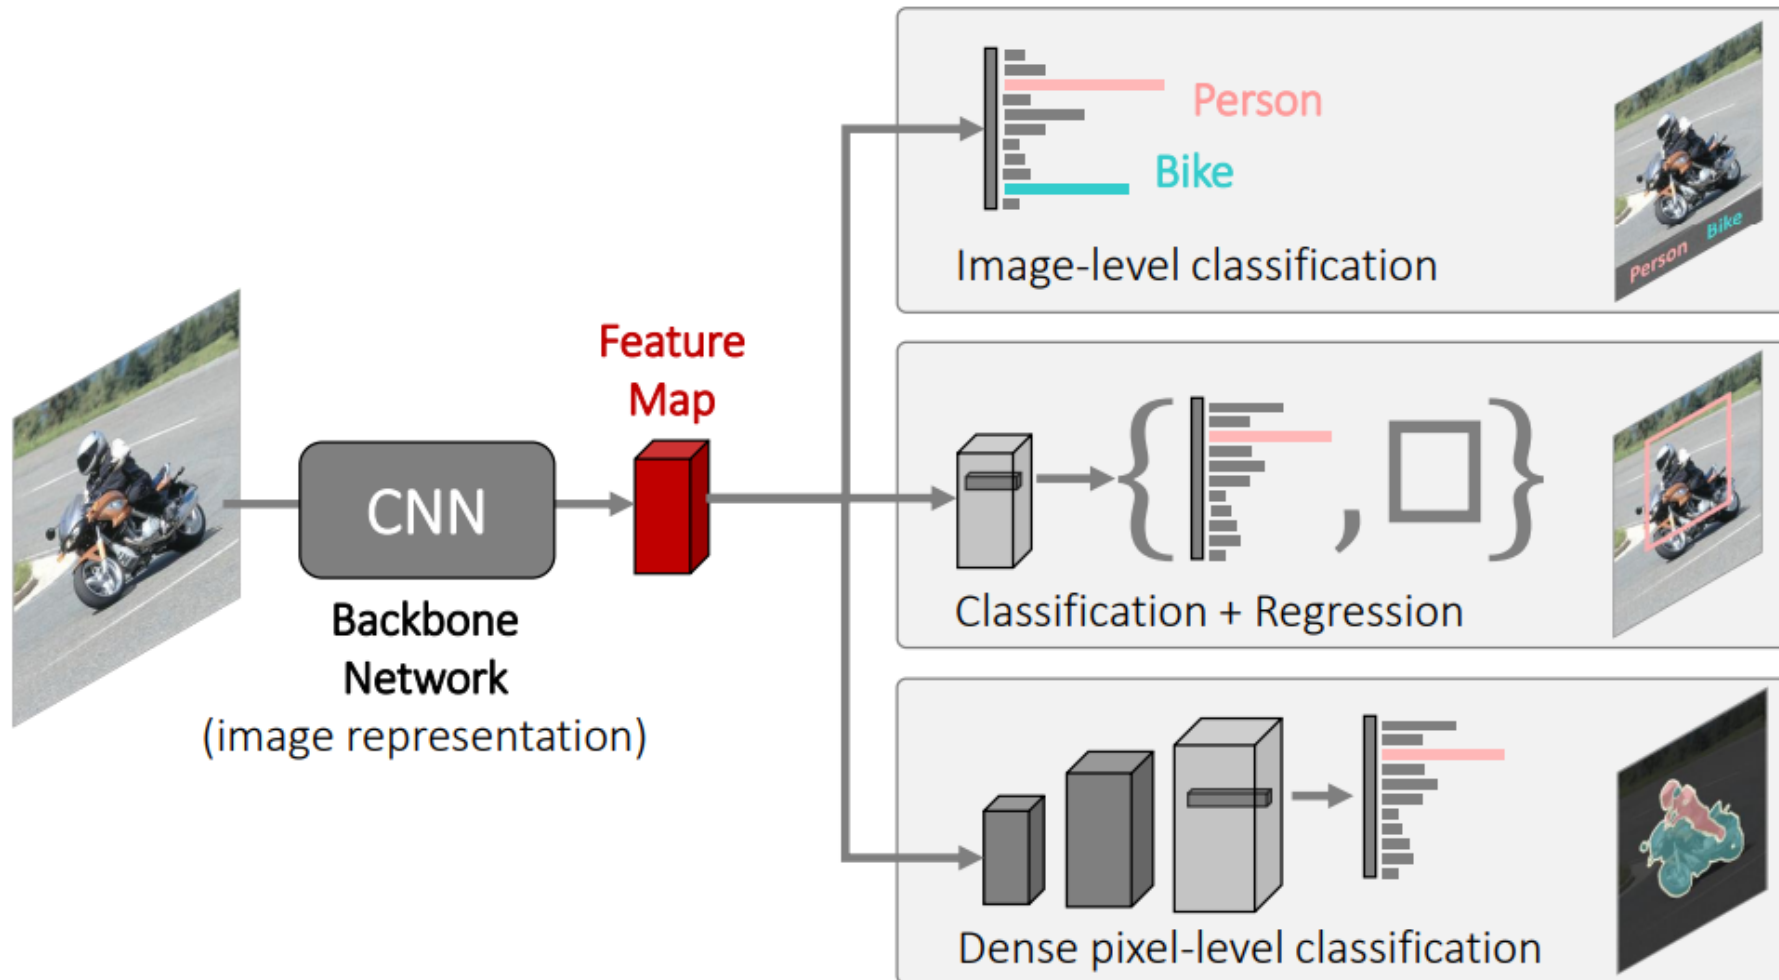
→ **Cat**

Goal : Classify the label for each image

Classification task




Classification CNN: Backbone Network



CNN Architectures

- AlexNet
- VGG
- GoogLeNet
- ResNet
- ...

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners




AlexNet (2012)

- Successful CNN image classification model
 - Based on LeNet5 CNN design (Yann Lecun et al, 1989)
 - **Computationally expensive, but feasible due to GPUs**
 - Parallel computation
 - Winner of ILSVRC-2012 competition by a large margin,
 - ImageNet Large-Scale Visual Recognition Challenge
 - a top-5 error of 15.3%, more than 10.8 percentage points lower than that of the runner up.
 - Transfer to significant gains in a variety of domains

LeNet-5 (1998)

- A very simple CNN architecture by Yann LeCun in 1998

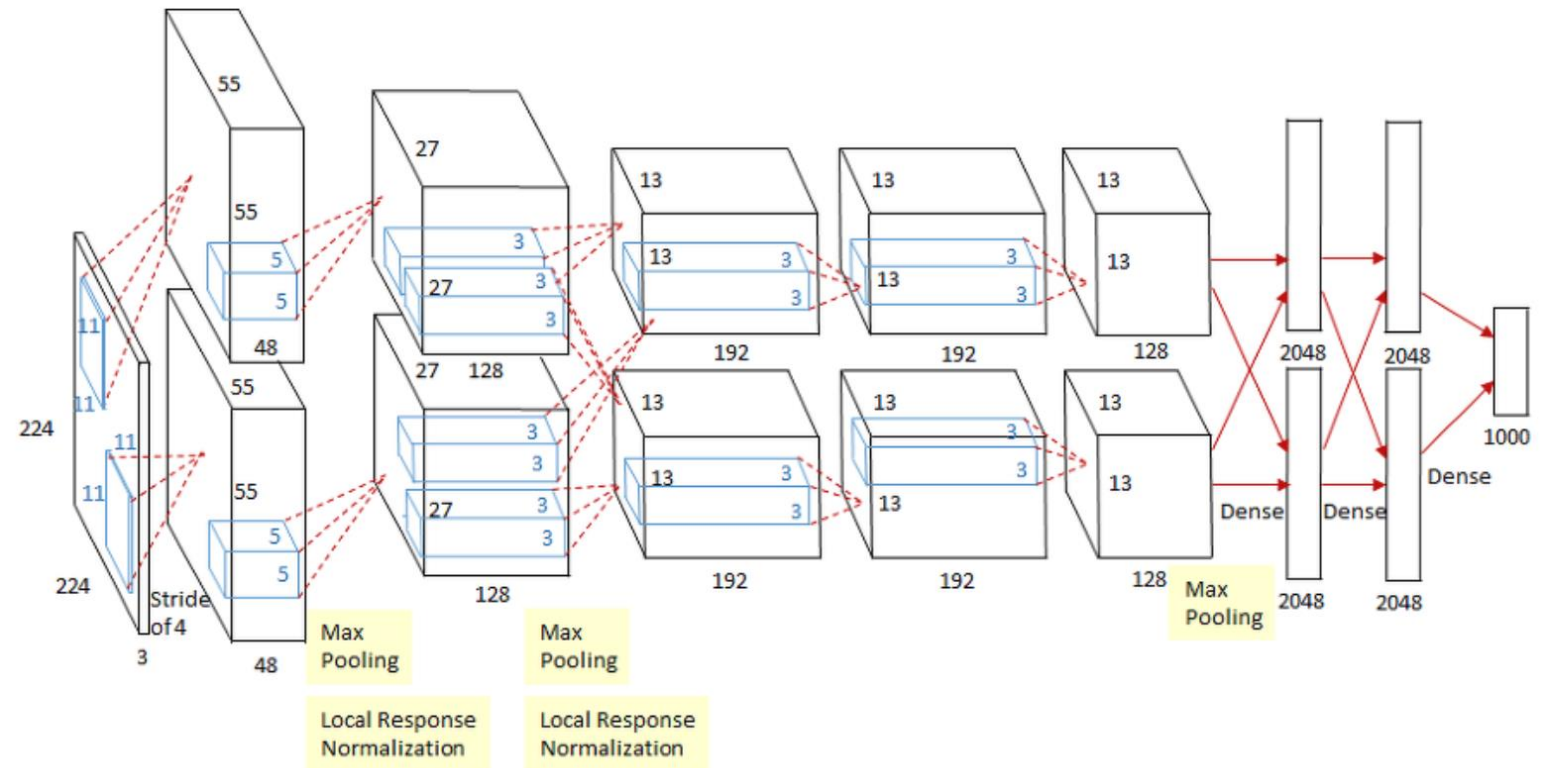
Overall architecture: Conv – Pool – Conv – Pool – FC – FC



[Lecun89] Y. LeCun et al.: **Handwritten Digit Recognition with a Back-Propagation Network**. NIPS 1989

AlexNet

- Architecture
 - Conv1
 - Maxpool1
 - Norm1
 - Conv2
 - Maxpool2
 - Norm2
 - Conv3
 - Conv4
 - Conv5
 - Maxpool3
 - FC6
 - FC7
 - FC8



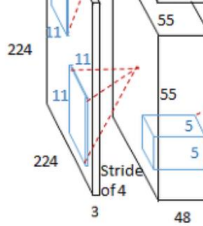
Overall architecture:

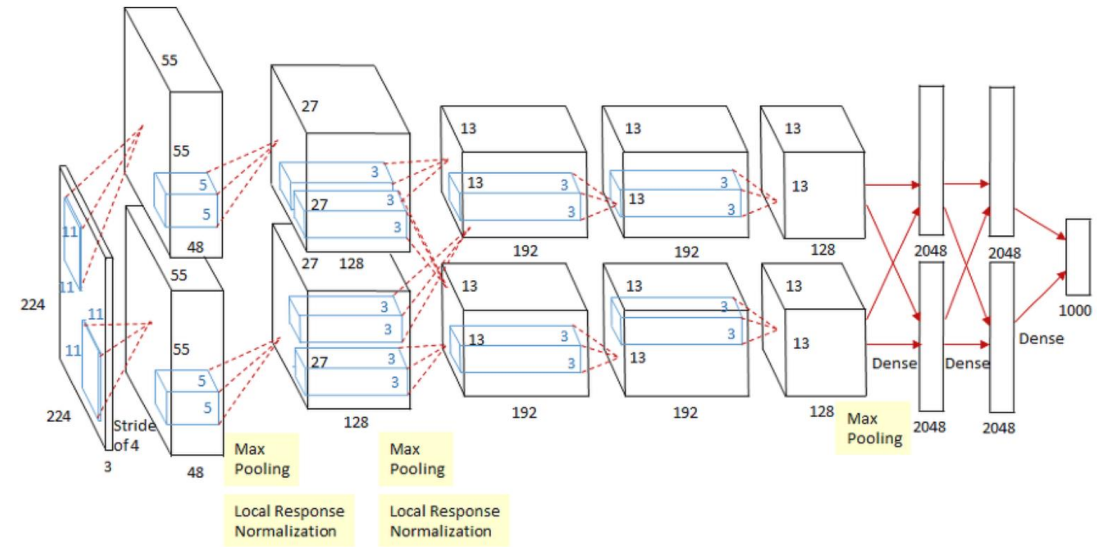
Conv – Pool – LRN – Conv – Pool – LRN – Conv – Conv – Conv – Pool – FC – FC – FC

LeNet-5 vs AlexNet

- Compared to LeNet-5, AlexNet is ...
 - Bigger (more layers, more parameters)
 - Trained with a larger amount of data (ImageNet)
 - Use different activation function & regularization techniques (dropout)
 - Local Response Normalization
 - Parallel computation via GPUs

AlexNet

- Input: **227 * 227 * 3** images
 - After Conv1 (96 11*11 filters, stride 4): **55 * 55 * 96**
 - After Pool1 (3*3 filters, stride 2): **27*27*96**
 - ...
 - Output size of convolution = $\frac{\text{input size} - \text{filter size} + (2 \times \text{padding})}{\text{stride}} + 1$
- 



AlexNet

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

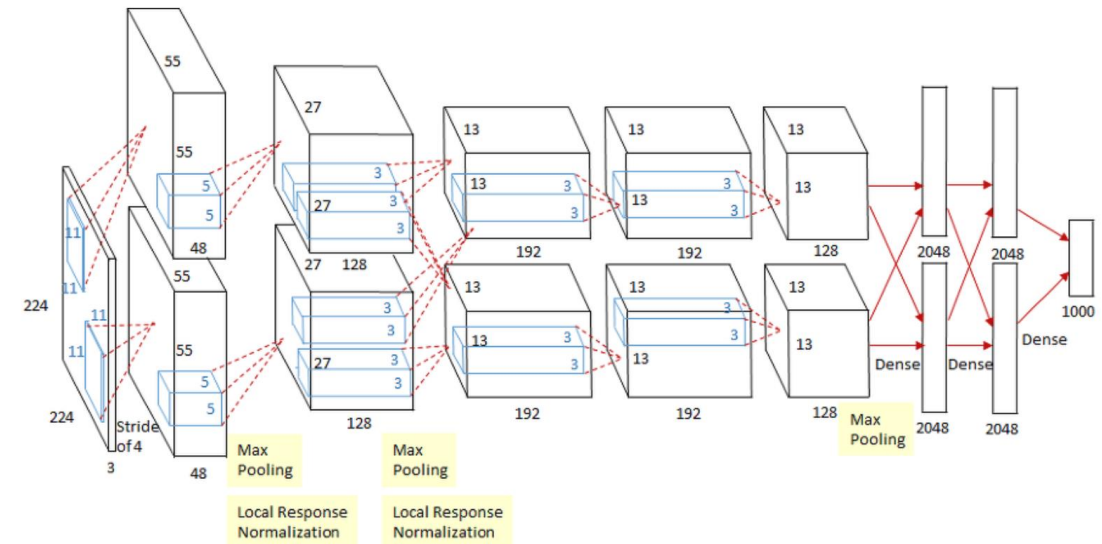
[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



AlexNet

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

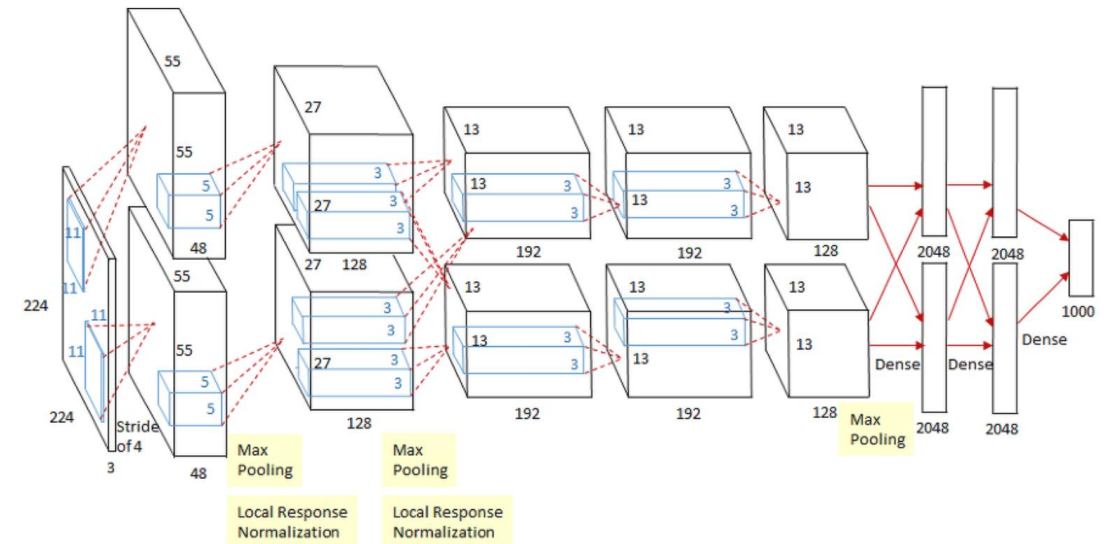
[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons


[1000] **FC8**: 1000 neurons (class scores)



Other Details:

- First use of ReLU
- Used Local Response Normalization
- Heavy data augmentation
- Dropout 0.5
- Batch size 128
- Learning rate 1e-2, reduced by 10 when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% -> 15.4 %

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



VGG network (2014)

- Proposed by Oxford VGG team in 2014 ILSVRC (2nd rank)
- Comparison
 - More deeper network than AlexNet (Alexnet 8 layers -> 16 – 19 layers)
 - More simple structure than GoogleNet (1st rank)
- Filter
 - Only **3 x 3** filter & 1 x 1 filter
 - Less parameter for same receptive field -> Regularization
- 11.7% top 5 error in ILSVRC'13 -> 7.3% top 5 error in ILSVRC'14


VGG network

- Why use smaller filters? (3×3 conv)
 - Stack of three 3×3 conv (stride 1) layers has the same receptive field as one 7×7 conv layer
 - Receptive field: The region in the input space that a particular CNN's feature is looking at

VGG network

- Why use smaller filters? (3*3 conv)
 - Stack of three 3*3 conv (stride 1) layers has the same receptive field as one 7*7 conv layer
 - Receptive field: The region in the input space that a particular CNN's feature is looking at


- ex)



VGG network

- Why use smaller filters? (3*3 conv)
 - Stack of three 3*3 conv (stride 1) layers has the same receptive field as one 7*7 conv layer
 - Receptive field: The region in the input space that a particular CNN's feature is looking at


- ex)



VGG network

- Why use smaller filters? (3×3 conv)
 - Stack of three 3×3 conv (stride 1) layers has the same receptive field as one 7×7 conv layer
 - Receptive field: The region in the input space that a particular CNN's feature is looking at

- ex)




VGG network

- Why use smaller filters? (3*3 conv)
 - Stack of three 3*3 conv (stride 1) layers has the same receptive field as one 7*7 conv layer
 - Receptive field: The region in the input space that a particular CNN's feature is looking at
- But we have deeper, more non-linearities
- And fewer parameters:
 - 3-(3*3) filters = 27
 - 1-(7*7) filter = 49


VGG16 network

- Configuration
 - Image size : $224 * 224 * 3$
 - All conv layer: Stride 1, Padding 1, filter $3*3$ or $1*1$
 - Maintain resolution (H * W size) of feature
 - Maxpooling: $2 * 2$ window, 2 stride




VGG16 network

Input: [224*224*3]
Conv3-64: [224*224*64]
Conv3-64: [224*224*64]
Pool2: [112*112*64]
Conv3-128: [112*112*128]
Conv3-128: [112*112*128]
Pool2: [56*56*128]
Conv3-256: [56*56*256]
Conv3-256: [56*56*256]
Conv3-256: [56*56*256]
Pool2: [28*28*256]
Conv3-512: [28*28*512]
Conv3-512: [28*28*512]
Conv3-512: [28*28*512]
Pool2: [14*14*512]
Conv3-512: [14*14*512]
Conv3-512: [14*14*512]
Conv3-512: [14*14*512]
Pool2: [7*7*512]
FC: [1*1*4096]
FC: [1*1*4096]
FC: [1*1*1000]



VGG16 network

- Other details
 - ILSVRC'14 2nd winner
 - No Local Response Normalization
 - Use VGG16 or VGG19 (VGG19 slightly better)
 - Use ensembles for best results



VGG network

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Pytorch code pipeline

- Define Dataset -> 학습하고 싶은 데이터셋을 정의
- Define DataLoader -> 데이터셋의 배치화
- Define model -> 학습할 모델 정의
- Define loss (criterion) -> loss 계산
- Define optimizer -> gradient descent 수행

-> Start Training and Testing

Pytorch code example

- **Define Dataset** →
- Define DataLoader →
- Define model →
- Define loss (criterion) →
- Define optimizer →
- Training and Testing →

```
from torchvision.datasets import MNIST
import torchvision.transforms as transforms

train_data = MNIST('./data/train', train=True, download=True, transform=transforms.ToTensor())
test_data = MNIST('./data/test', train=False, download=True, transform=transforms.ToTensor())
```

Pytorch code example

- Define Dataset →
- **Define DataLoader** →
- Define model →
- Define loss (criterion) →
- Define optimizer →
- Training and Testing →

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(train_data, batch_size=16, shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=16, shuffle=False)
```

Pytorch code example

- Define Dataset →
- Define DataLoader →
- **Define model** →
- Define loss (criterion) →
- Define optimizer →
- Training and Testing →

```
from torch import nn

class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_layer = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_layer(x)
        return logits

model = Model()
```

Pytorch code example

- Define Dataset →
- Define DataLoader →
- Define model →
- **Define loss (criterion) →**
- **Define optimizer →**
- Training and Testing →

```
import torch.optim as optim

optimizer = optim.Adam(model.parameters(), lr=0.01, weight_decay=0.001)
loss_function = nn.MSELoss()
```

Pytorch code example


- Define Dataset →
- Define DataLoader →
- Define model →
- Define loss (criterion) →
- Define optimizer →
- **Training and Testing** →

```
model.train()
for x, y in train_dataloader:
    prediction = model(x)
    loss = loss_function(prediction, y)


    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

```
model.eval()
for x, y in test_dataloader:
    prediction = nn.Softmax(dim=1)(model(x))
    y_pred = prediction.argmax(1)
    print("Predicted class: {y_pred}, True label: {y}")
```


Ex) MNIST dataset



- Hand-written number dataset
- Total 10 class (0 ~ 9)
- 28X28 size, grayscale
- Training sample : 60,000
- Test sample : 10,000



Ex) CIFAR-10 dataset

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



- 10-class dataset
- 32 x 32 size, RGB image
- Training sample : 50,000
- Test sample : 10,000

Softmax function

- Function that takes as input a vector of K real numbers, and normalizes it into a probability distribution

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad \text{for } i = 1, \dots, K \text{ and } \mathbf{z} = (z_1, \dots, z_K) \in \mathbb{R}^K$$

- Also known as, normalized exponential function

Cross entropy loss

- To measure the distance between probability p and q

$$H(P, Q) = - \sum_x P(x) \log Q(x)$$

- Ex) Given that ground truth label $P(x) = (1, 0)$
 - if $Q(x) = (0, 1)$, $-P(x) \log Q(x) = -[1 \ 0] \begin{bmatrix} \log 0 \\ \log 1 \end{bmatrix} = -(-\infty + 0) = \infty$
 - if $Q(x) = (1, 0)$, $-P(x) \log Q(x) = -[1 \ 0] \begin{bmatrix} \log 1 \\ \log 0 \end{bmatrix} = -(0 + 0) = 0$

Datasets and DataLoader

- All datasets are subclass of `torch.utils.data.Dataset`
 - You can make your own datasets to inherit Dataset class
 - Some famous datasets is provided in `torchvision.datasets`
 - CIFAR-10, MNIST, etc
- `DataLoader` reads datasets and make the batch

torch.utils.data.DataLoader

- It reads datasets and make batch for training and inference
- Arguments
 - dataset – target dataset (torch.utils.data.Dataset)
 - batch_size – batch size for training or inference
 - shuffle – if true, dataloader shuffle the data in datasets randomly
 - sampler, batch_sampler, num_workers, pin_memorys
- for batch_idx, (data, label) in enumerate(trainloader)

Exercise 1.

Training a classifier on CIFAR-10

- Step)
 - Load and normalizing the CIFAR 10 training and test datasets using torchvision
 - Visualize first 8 images from the train dataset
 - **1. Define a convolutional Neural Network**
 - **2. Define a loss function and optimizer (CrossEntropy, SGD)**
 - **3. Train the network on the training data**
 - Test the network on the test data

Visualize first 8 images from the train dataset


```
import matplotlib.pyplot as plt
from torchvision.transforms.functional import to_pil_image
```



import library for plotting/visualizing data

```
plt.figure(figsize=(16,4))
for i in range(8):
    img = to_pil_image(train_data[i][0]/2 + 0.5)
    plt.subplot(1,8,i+1)
    plt.imshow(img)
plt.show()
```

```
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
print("Labels: ", end="")
for i in range(8): print(classes[train_data[i][1]], end=' | ')
```



Labels: frog | truck | truck | deer | car | car | bird | horse |

Visualize first 8 images from the train dataset


```
import matplotlib.pyplot as plt
from torchvision.transforms.functional import to_pil_image
```

```
plt.figure(figsize=(16,4))
for i in range(8):
    img = to_pil_image(train_data[i][0]/2 + 0.5)
    plt.subplot(1,8,i+1)
    plt.imshow(img)
plt.show()
```



Define figure and plot images

```
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
print("Labels: ", end="")
for i in range(8): print(classes[train_data[i][1]], end=' | ')
```



Labels: frog | truck | truck | deer | car | car | bird | horse |

Visualize first 8 images from the train dataset


```
import matplotlib.pyplot as plt
from torchvision.transforms.functional import to_pil_image
```

```
plt.figure(figsize=(16,4))
for i in range(8):
    img = to_pil_image(train_data[i][0]/2 + 0.5)
    plt.subplot(1,8,i+1)
    plt.imshow(img)
plt.show()
```

```
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
print("Labels: ", end="")
for i in range(8): print(classes[train_data[i][1]], end=' | ')
```



Print labels corresponding to the images



Labels: frog | truck | truck | deer | car | car | bird | horse |

- Load and normalizing the CIFAR 10 training and test datasets using torchvision

```
from torchvision.datasets import CIFAR10
import torchvision.transforms as transforms

transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

train_data = CIFAR10("./data", train=True, download=True, transform=transform)
test_data = CIFAR10("./data", train=False, download=True, transform=transform)

print(len(train_data))
print(len(test_data))
```

Files already downloaded and verified

Files already downloaded and verified

50000

10000

- Load and normalizing the CIFAR 10 training and test datasets using torchvision

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(train_data, batch_size=8, shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=8, shuffle=False)
```

1. Define a convolutional Neural Network

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 8, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2,2)
        )
        self.layer2 = nn.Sequential(
            nn.Conv2d(8, 16, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2,2)
        )

        self.fc_layer = nn.Linear(16 * 8 * 8, 10)

    def forward(self, x):
        x1 = self.layer1(x)
        x2 = self.layer2(x1)

        x2 = torch.flatten(x2, 1)
        x3 = self.fc_layer(x2)
        out = F.softmax(x3, dim=1)
        return out

net = CNN()
```

```
if torch.cuda.is_available(): device = torch.device('cuda')
else: device = torch.device('cpu')

net = net.to(device)
```

✓ 0.3s

2. Define a loss function and optimizer (CrossEntropy, SGD)

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

3. Train the network on the training data

```
for epoch in range(4): # 데이터셋을 수차례 반복합니다.

    running_loss = 0.0
    net.train()
    for i, data in enumerate(train_dataloader, 0):
        # [inputs, labels]의 목록인 data로부터 입력을 받은 후;
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        # 변화도(Gradient) 매개변수를 0으로 만들고
        optimizer.zero_grad()

        # 순전파 + 역전파 + 최적화를 한 후
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # 통계를 출력합니다.
        running_loss += loss.item()
    if i % 2000 == 1999: # print every 2000 mini-batches
        print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.3f}', end=' | ')
        running_loss = 0.0
```

Test the network on the test data

```
correct = 0
total = 0
net.eval()
# 학습 중이 아니므로, 출력에 대한 변화도를 계산할 필요가 없습니다
with torch.no_grad():
    for data in test_dataloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)

        # 신경망에 이미지를 통과시켜 출력을 계산합니다
        outputs = net(images)
        # 가장 높은 값(energy)을 갖는 분류(class)를 정답으로 선택하겠습니다
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy on the 10000 test images: {100 * correct // total} %')

print('Finished Training')
```

```
[1, 2000] loss: 2.279 | [1, 4000] loss: 2.187 | [1, 6000] loss: 2.149 | Accuracy on the 10000 test images: 33 %
[2, 2000] loss: 2.122 | [2, 4000] loss: 2.101 | [2, 6000] loss: 2.072 | Accuracy on the 10000 test images: 40 %
[3, 2000] loss: 2.050 | [3, 4000] loss: 2.036 | [3, 6000] loss: 2.028 | Accuracy on the 10000 test images: 44 %
[4, 2000] loss: 2.014 | [4, 4000] loss: 1.999 | [4, 6000] loss: 2.002 | Accuracy on the 10000 test images: 47 %
Finished Training
```


Exercise 2.

Train VGG (A) Network on CIFAR-10

Step:

1. Implement VGG11 network
2. Change the fully connected part (**use only one fc**)

- Conv – input channels: 3, output channels: 64
- Maxpooling
- Conv – input channels: 64, output channels: 128
- Maxpooling
- Conv – input channels: 128, output channels: 256
- Conv – input channels: 256, output channels: 256
- Maxpooling
- Conv – input channels: 256, output channels: 512
- Conv – input channels: 512, output channels: 512
- Maxpooling
- Conv – input channels: 512, output channels: 512
- Conv – input channels: 512, output channels: 512
- Maxpooling

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Every conv layer: kernel size: 3, padding: 1
Every max pooling layer: kernel size 2, stride 2

Exercise 2-2.

Train VGG (D) Network on CIFAR-10

Step:

1. Implement VGG16 network
2. Change the fully connected part (**use three fc layers**)
3. Apply dropout of 0.5 to fc layers

- Conv – input channels: 3, output channels: 64
- Conv – input channels: 64, output channels: 64
- Maxpooling
- Conv – input channels: 64, output channels: 128
- Conv – input channels: 128, output channels: 128
- Maxpooling
- Conv – input channels: 128, output channels: 256
- Conv – input channels: 256, output channels: 256
- Conv – input channels: 256, output channels: 256
- Maxpooling
- Conv – input channels: 256, output channels: 512
- Conv – input channels: 512, output channels: 512
- Conv – input channels: 512, output channels: 512
- Maxpooling
- Conv – input channels: 512, output channels: 512
- Conv – input channels: 512, output channels: 512
- Conv – input channels: 512, output channels: 512
- Maxpooling

Every conv layer: kernel size: 3, padding: 1
Every max pooling layer: kernel size 2, stride 2

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-12 conv3-12	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-25 conv3-25 conv1-25	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-51 conv3-51 conv1-51	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-51 conv3-51 conv1-51	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					