

Classification (2)

POSTECH MIP Lab.

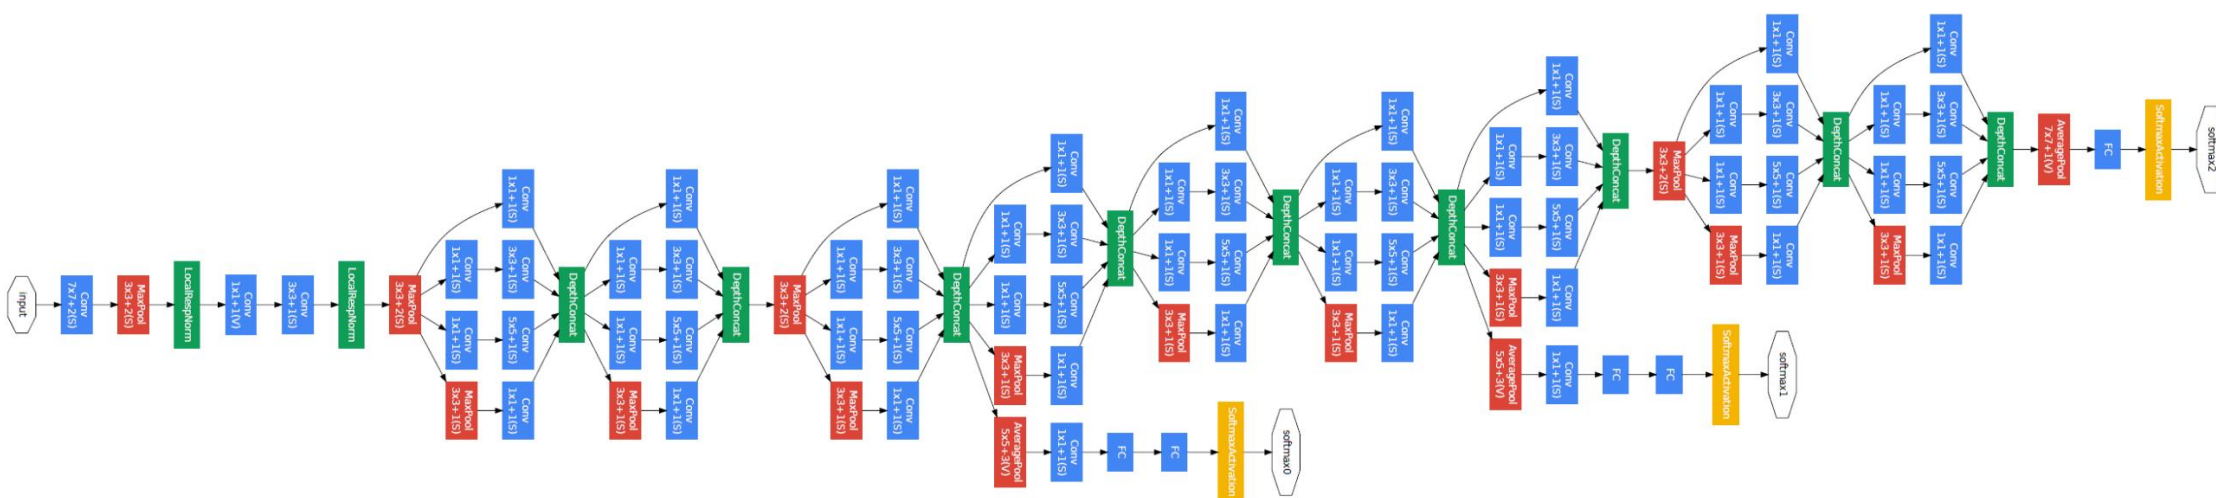
TA: Joonhyuk Park, Seunghun Baek, Soojin Hwang

Goal

- Understand classification task
- Understand the development of CNN based classification models (AlexNet, VGG, **GoogLeNet**, **ResNet**)
- Learn how to implement CNN models from architecture

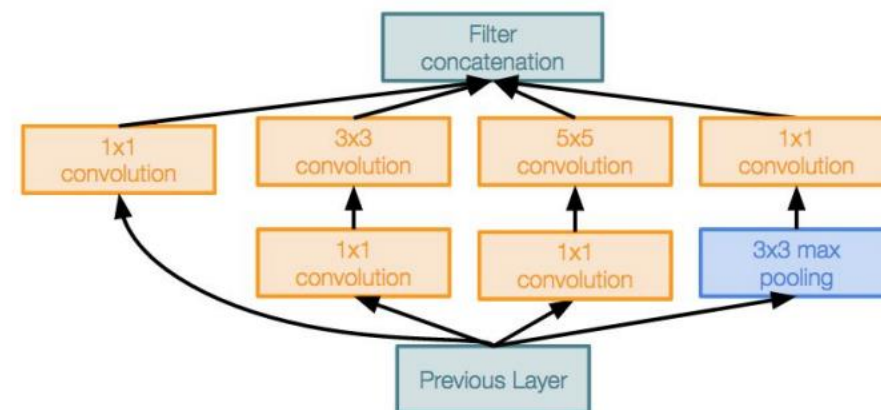
GoogLeNet (2014)

- Deeper networks (22 layers)
- Use efficient Inception module
 - Only 5 million parameters (12 times less than AlexNet)
- ILSVRC'14 winner



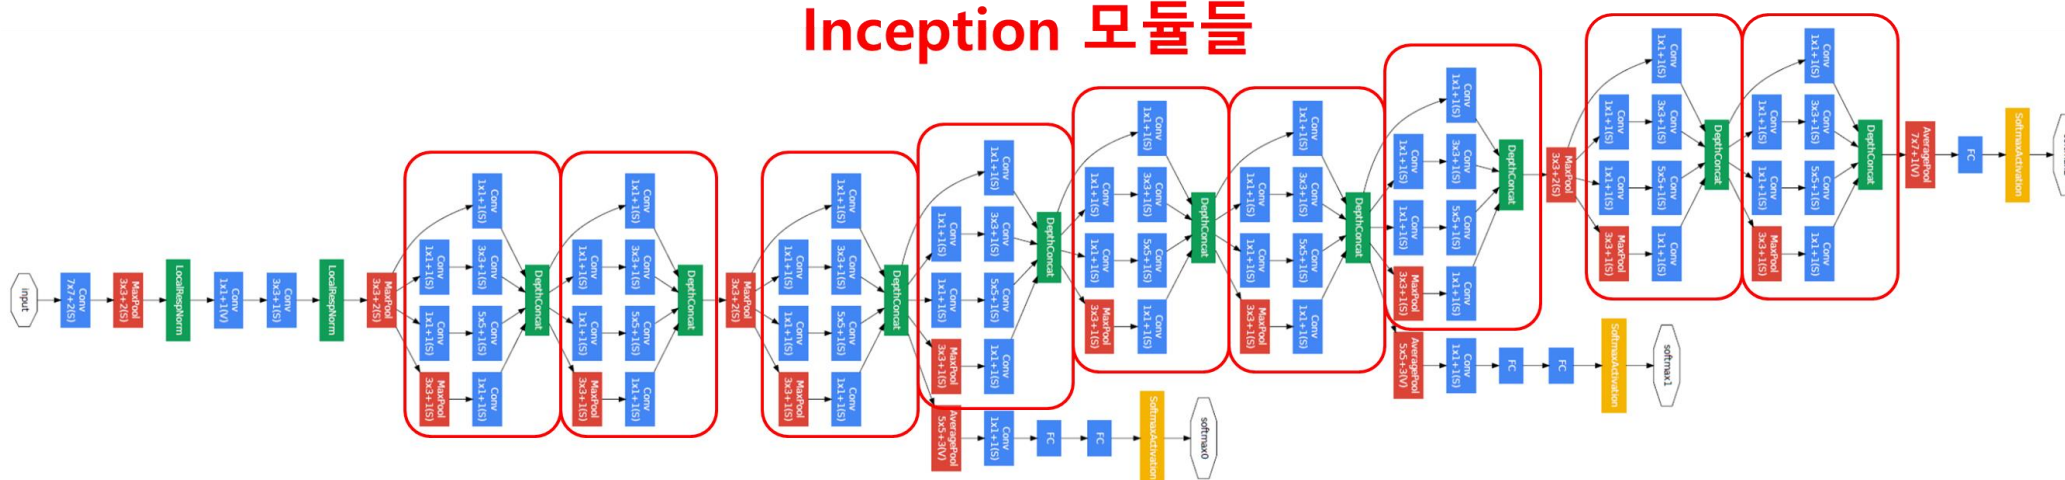
GoogLeNet (2014)

- Inception Module?



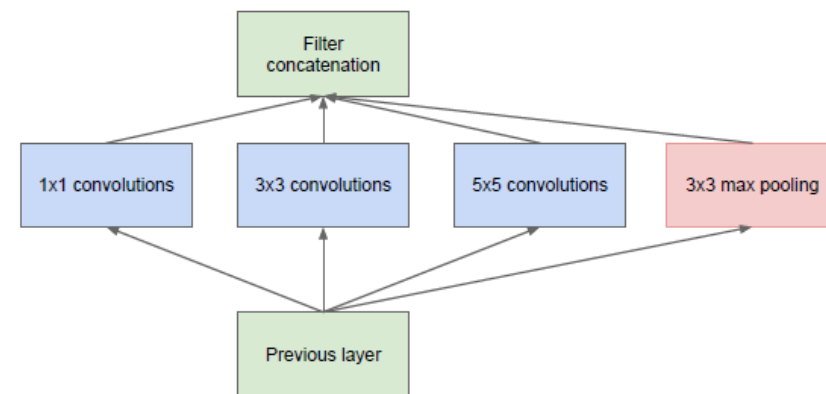
Inception module

Inception 모듈들

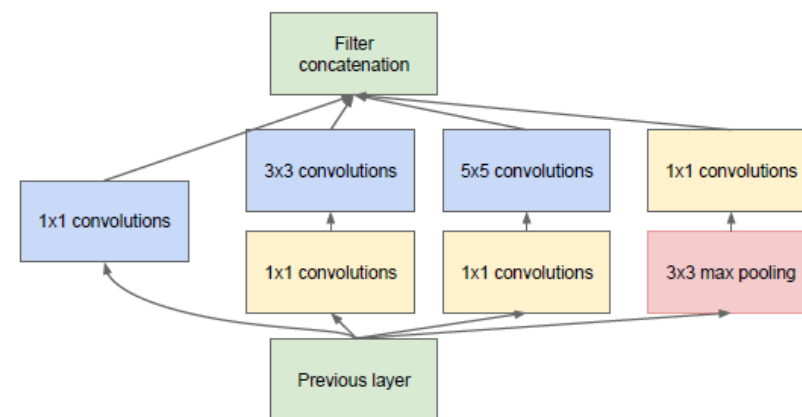


GoogLeNet (2014)

- Inception Module?
 - Apply parallel filter operations on the input from previous layer
 - 1*1 conv, 3*3 conv, 5*5 conv, 3*3 maxpool
 - -> problem: computational complexity



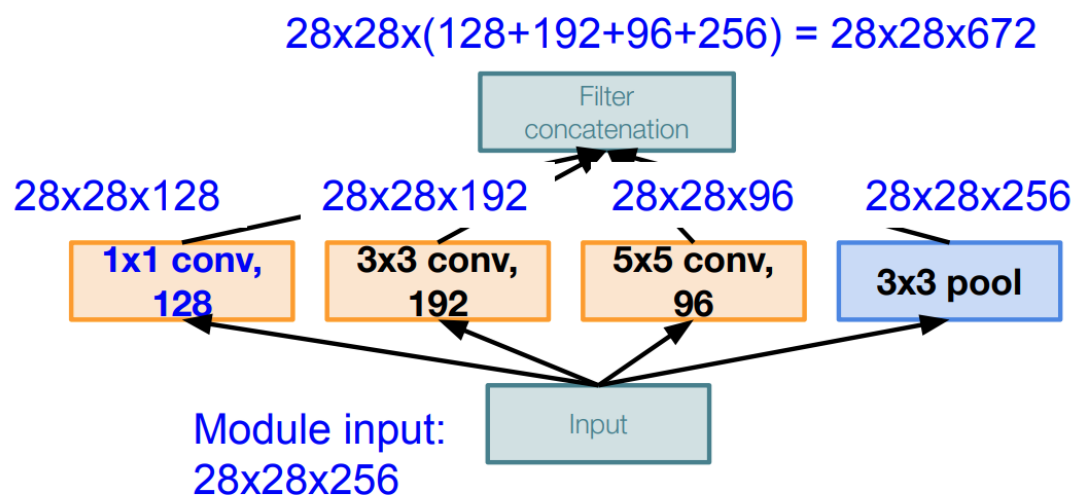
(a) Inception module, naïve version



(b) Inception module with dimensionality reduction

GoogLeNet (2014)

- Inception Module?
 - problem: computational complexity
 - Ex)



Naive Inception module

Conv Ops:

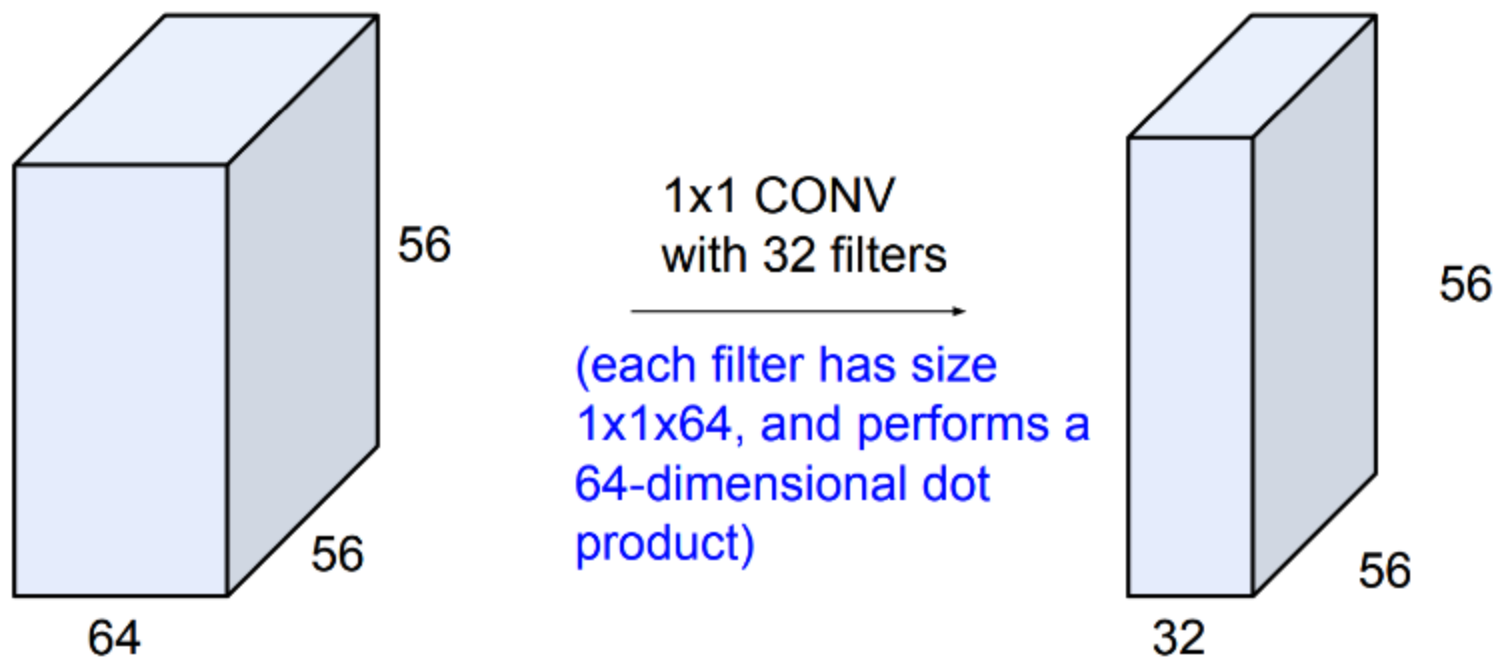
[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$
[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$
[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

Total: 854M ops

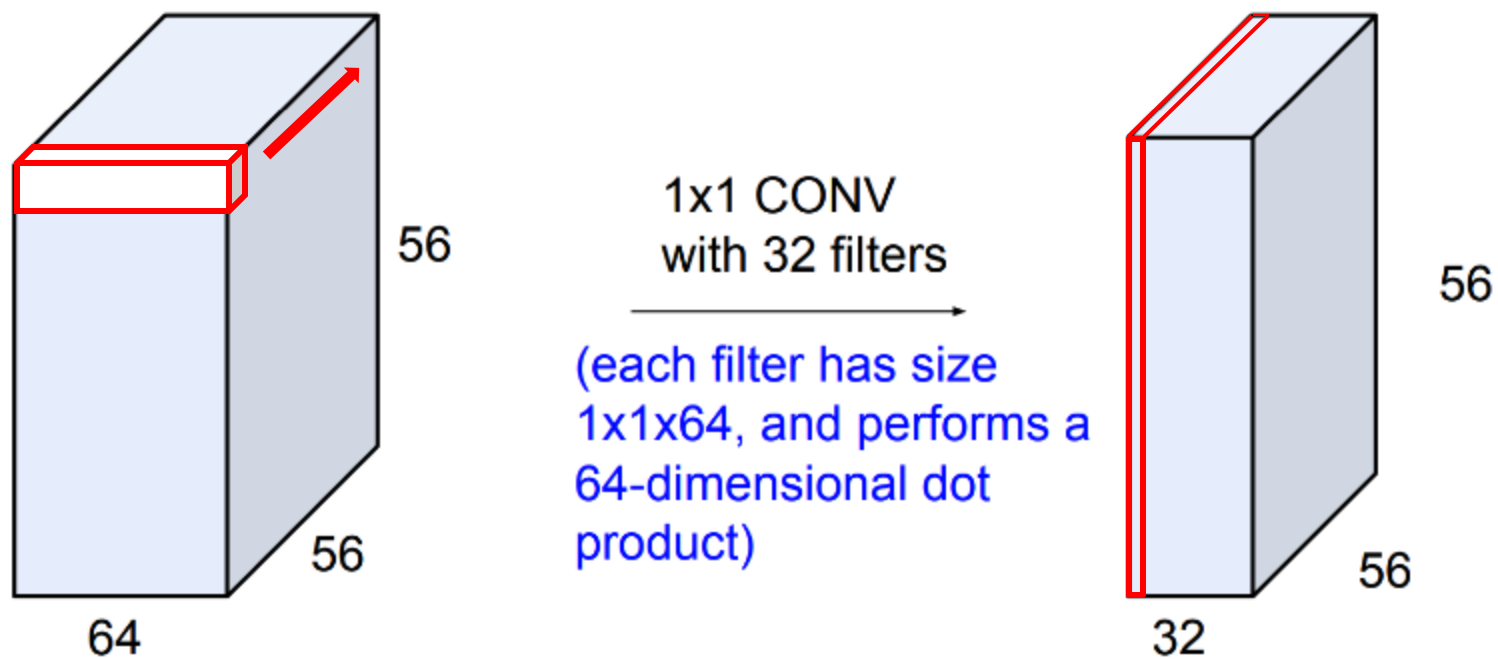
GoogLeNet (2014)

- Inception Module?
 - problem: computational complexity
 - Solution: bottleneck layer (1×1 convolutions)

1 x 1 Convolution



1 x 1 Convolution

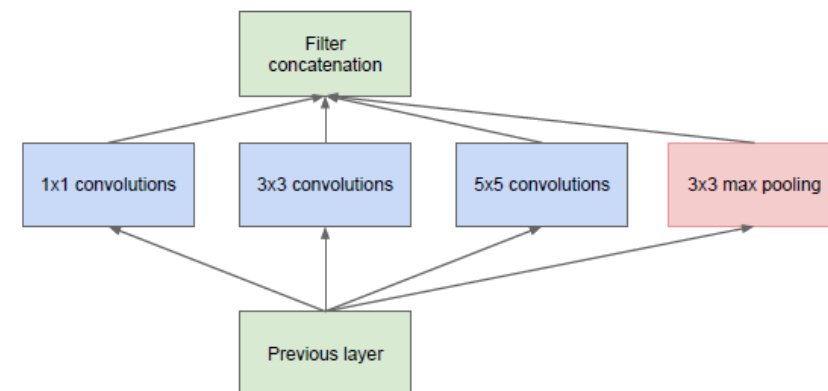


1 x 1 Convolution

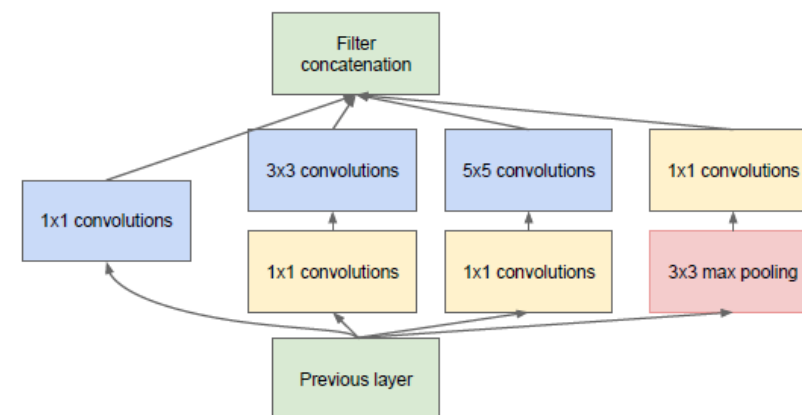
- Preserves spatial dimensions, but reduces depth!
- Projects depth to lower dimension (combination of feature maps)

GoogLeNet (2014)

- Inception Module?
 - problem: computational complexity
 - Solution: bottleneck layer (1*1 convolutions)



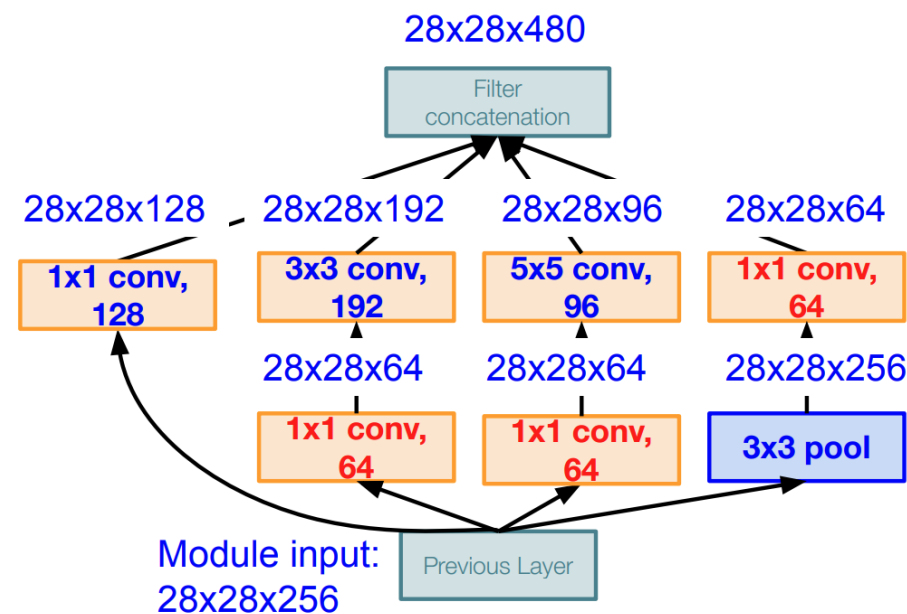
(a) Inception module, naïve version



(b) Inception module with dimensionality reduction

GoogLeNet (2014)

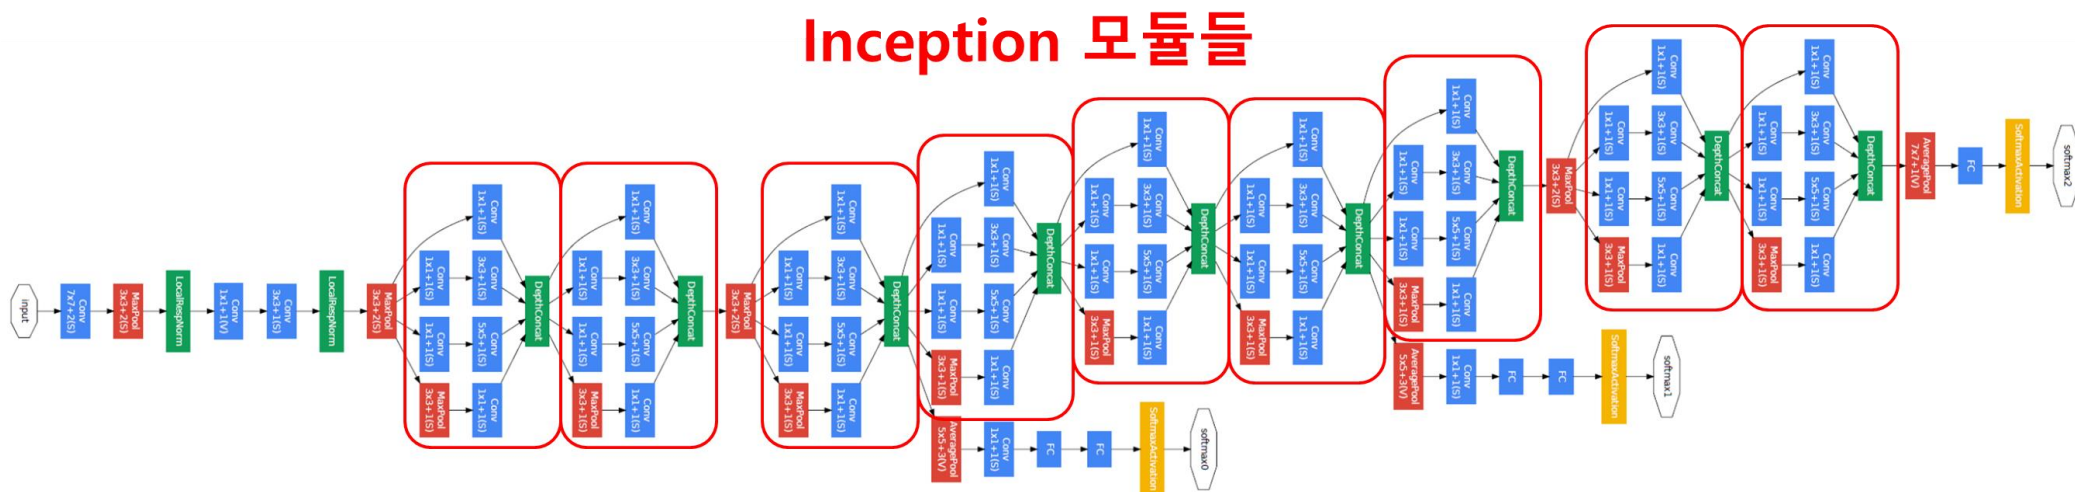
- Inception Module?
 - problem: computational complexity
 - Solution: bottleneck layer (1*1 convolutions)




Inception module with dimension reduction

GoogLeNet (2014)

- Architecture details
 - Stack of inception modules
 - Use of auxiliary classifiers
 - No FC layer, use global average pooling



Common trend : Revolution of depth





Residual Network (2015)

- Deeper networks with residual connection (152-layers)
- ILSVRC'15 winner

Residual Network (2015)

- Deep network has some problems
 - Vanishing/Exploding gradient
 - Too much parameter




- Residual Network solves these problems


Residual Network (2015)

- Residual Network benefits
 - Easy to learn the residual
 - Solve gradient vanishing problem (미분해도 항상 1 이상)

- **Plaint net**

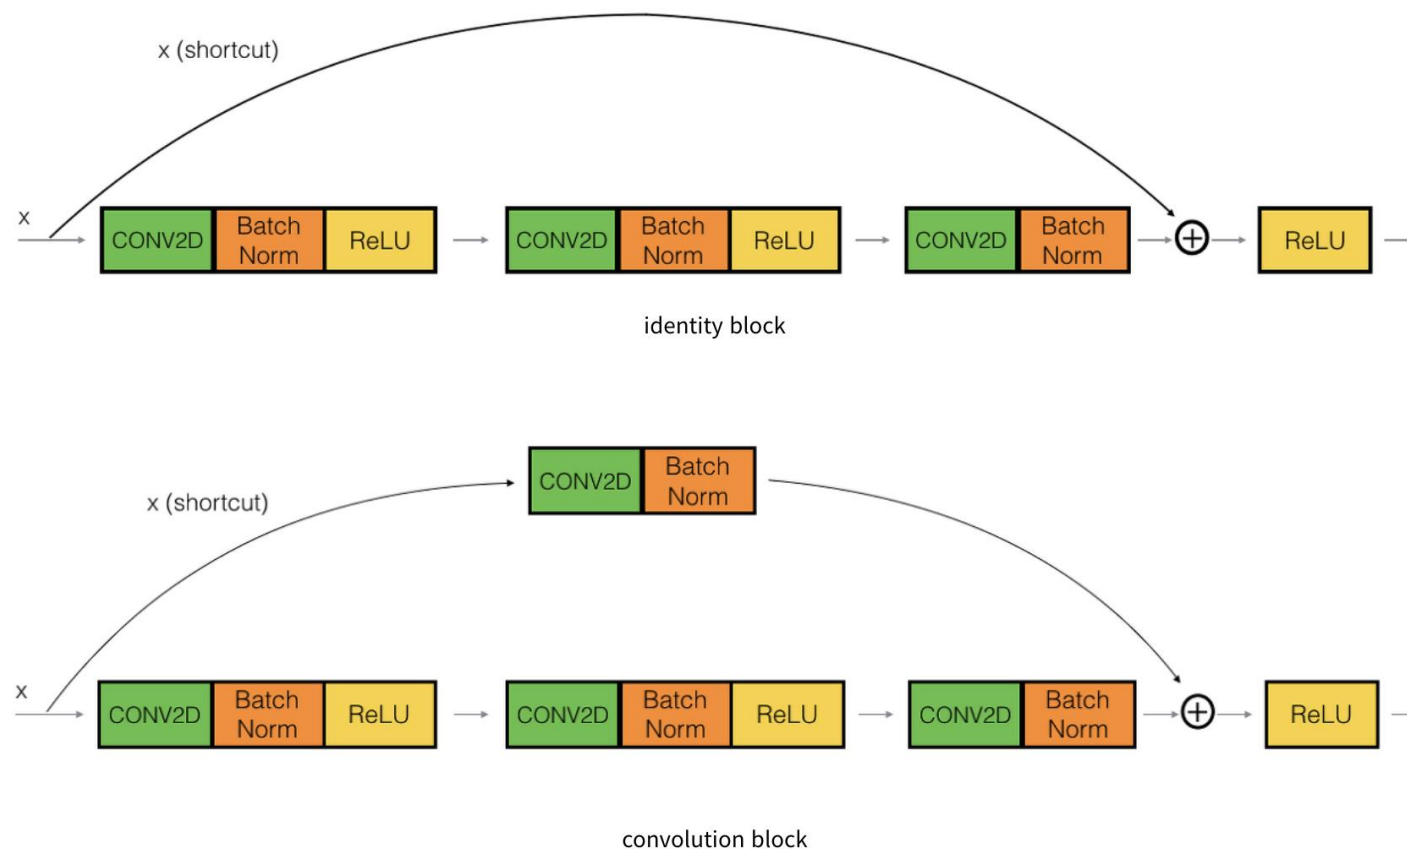


- **Residual net**



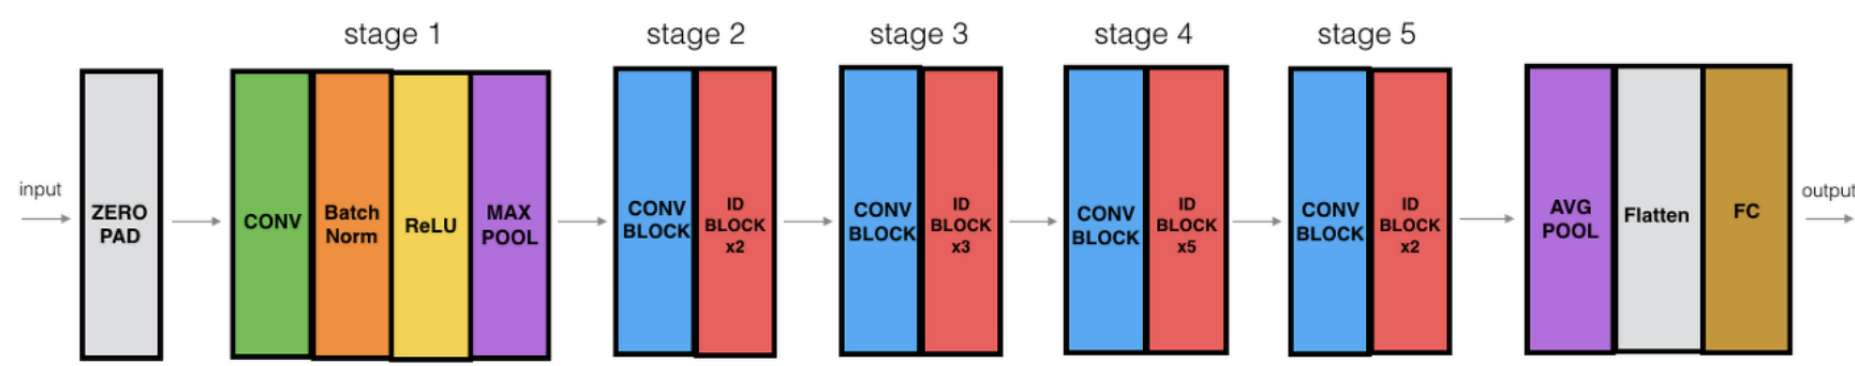
Residual Network (2015)

- Architecture

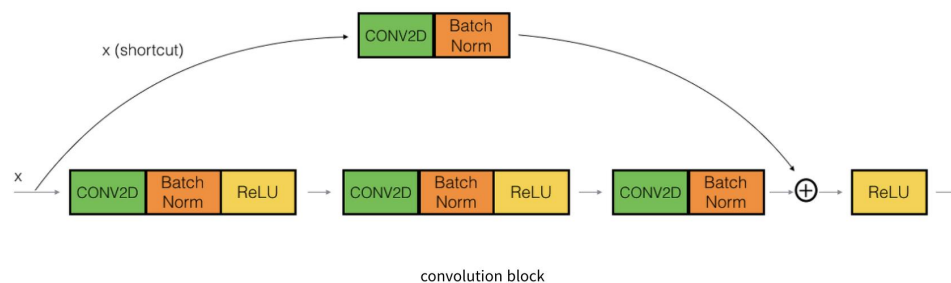



Residual Network (2015)

- Architecture

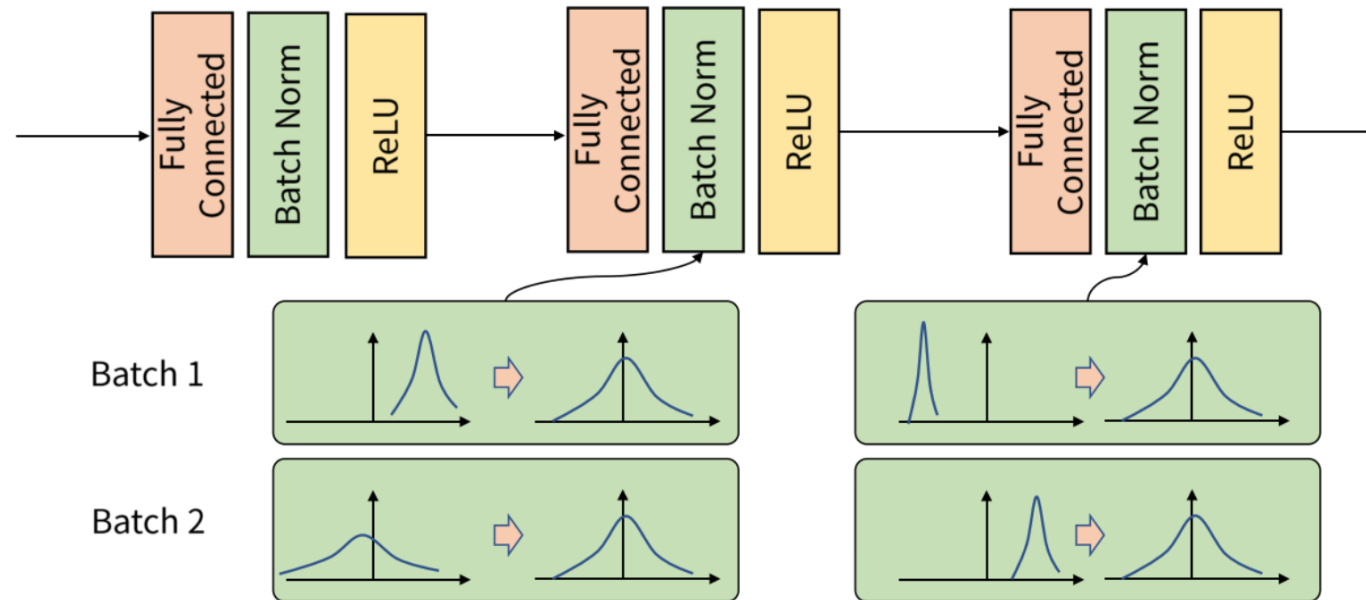


ResNet structure



Batch Normalization

- Internal Covariant Shift
 - Normalize each batch using mean and s.d of each batch



Residual Network (2015)

- Architecture


layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

Residual Network (2015)


method	top-5 err. (test)
VGG [41] (ILSVRC'14)	7.32
GoogLeNet [44] (ILSVRC'14)	6.66
VGG [41] (v5)	6.8
PReLU-net [13]	4.94
BN-inception [16]	4.82
ResNet (ILSVRC'15)	3.57

Exercise 3-1. Implement Residual Block


- Plain net




- Residual net




Error: 6.61%




Error: 8.17%



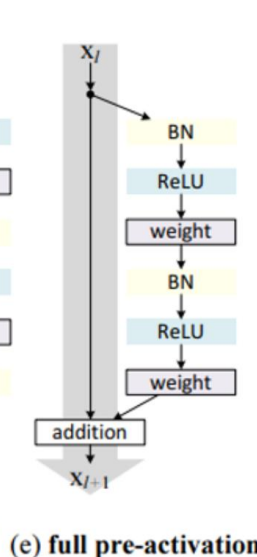
Error: 7.84%



Error: 6.71%

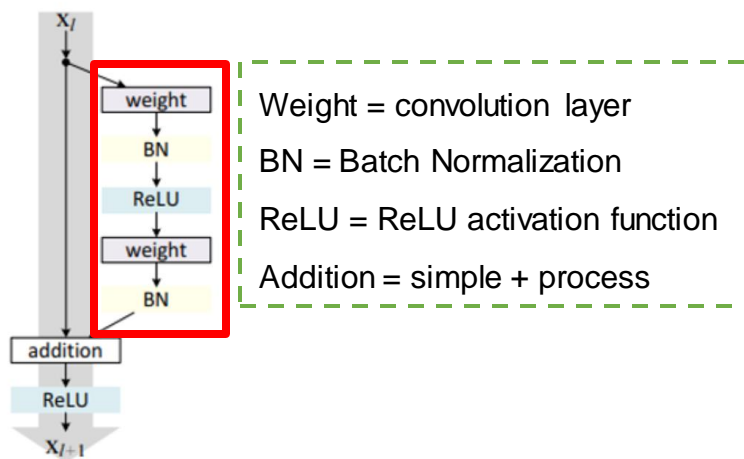


Error: 6.37%



Residual Block

Error: 6.61%



Weight = convolution layer

BN = Batch Normalization

ReLU = ReLU activation function

Addition = simple + process

(a) original

```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        # BatchNorm에 bias가 포함되어 있으므로, conv2d는 bias=False로 설정합니다.
        self.residual_function = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.shortcut = nn.Sequential()


        self.relu = nn.ReLU()

        # in_channels와 out_channels가 다를 경우 맞춰줌
        if in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(out_channels)
            )

    def forward(self, x):
        x = self.residual_function(x) + self.shortcut(x)
        x = self.relu(x)
        return x
```


Residual Block

Error: 6.61%



Weight = convolution layer
BN = Batch Normalization
ReLU = ReLU activation function
Addition = simple + process

(a) original

```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        # BatchNorm에 bias가 포함되어 있으므로, conv2d는 bias=False로 설정합니다.
        self.residual_function = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )


        self.shortcut = nn.Sequential()

        self.relu = nn.ReLU()

        # in_channels와 out_channels가 다를 경우 맞춰줌
        if in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(out_channels)
            )

    def forward(self, x):
        x = self.residual_function(x) + self.shortcut(x)
        x = self.relu(x)
        return x
```

Residual Block



```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        # BatchNorm에 bias가 포함되어 있으므로, conv2d는 bias=False로 설정합니다.
        self.residual_function = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.shortcut = nn.Sequential()


        self.relu = nn.ReLU()

        # in_channels와 out_channels가 다를 경우 맞춰줌
        if in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(out_channels)
            )

    def forward(self, x):
        x = self.residual_function(x) + self.shortcut(x)
        x = self.relu(x)
        return x
```

Residual Block

Error: 6.61%



(a) original

Weight = convolution layer
BN = Batch Normalization
ReLU = ReLU activation function
Addition = simple + process

```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()

        # BatchNorm에 bias가 포함되어 있으므로, conv2d는 bias=False로 설정합니다.
        self.residual_function = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.shortcut = nn.Sequential()

        self.relu = nn.ReLU()

        # in_channels와 out_channels가 다를 경우 맞춰줌
        if in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, bias=False),
                nn.BatchNorm2d(out_channels)
            )

    def forward(self, x):
        x = self.residual_function(x) + self.shortcut(x)
        x = self.relu(x)
        return x
```

Exercise 3-2. Implement ResNet-18

- Conv – input channels: 3, output channels: 64
- Maxpooling
- Conv – input channels: 64, output channels: 64 x 2
- Conv – input channels: 64, output channels: 64 x 2
- Conv – input channels: 64, output channels: 128
- Conv – input channels: 128, output channels: 128
- Conv – input channels: 128, output channels: 128 x 2
- Conv – input channels: 128, output channels: 256
- Conv – input channels: 256, output channels: 256
- Conv – input channels: 256, output channels: 256 x 2
- Conv – input channels: 256, output channels: 512
- Conv – input channels: 512, output channels: 512
- Conv – input channels: 512, output channels: 512 x 2
- Avgpooling
- FC

18-layer	34-layer	50-layer	101-layer	152-layer
7×7, 64, stride 2				
3×3 max pool, stride 2				
$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
average pool, 1000-d fc, softmax				
1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

Every conv layer: kernel size: 3, padding: 1
max pooling layer: kernel size 3, stride 2
Avg pooling layer: nn.AdaptiveAvgPool2d((1,1))

ResNet-18

18-layer	34-layer	50-layer	101-layer	152-layer
	7×7, 64, stride 2			
	3×3 max pool, stride 2			
$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	average pool, 1000-d fc, softmax			
1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

```

class ResNet18(nn.Module):
    def __init__(self, block, num_classes=10):
        super(ResNet18, self).__init__()

        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
        )

        self.conv2_1 = block(64, 64)
        self.conv2_2 = block(64, 64)

        self.conv3_1 = block(64, 128)
        self.conv3_2 = block(128, 128)

        self.conv4_1 = block(128, 256)
        self.conv4_2 = block(256, 256)

        self.conv5_1 = block(256, 512)
        self.conv5_2 = block(512, 512)

        self.avg_pool = nn.AdaptiveAvgPool2d((1,1)) # make output h*w = 1*1
        self.fc = nn.Linear(512, num_classes)

    def forward(self, x):
        # x: [batch, 3, 32, 32]

        x = self.conv1(x) # x: [batch, 64, 8, 8]

        x = self.conv2_1(x) # [batch, 64, 8, 8]
        x = self.conv2_2(x) # [batch, 64, 8, 8]

        x = self.conv3_1(x) # [batch, 128, 8, 8]
        x = self.conv3_2(x) # [batch, 128, 8, 8]

        x = self.conv4_1(x) # [batch, 256, 8, 8]
        x = self.conv4_2(x) # [batch, 256, 8, 8]

        x = self.conv5_1(x) # [batch, 512, 8, 8]
        x = self.conv5_2(x) # [batch, 512, 8, 8]

        x = self.avg_pool(x) # [batch, 512, 1, 1]
        x = x.view(x.size(0), -1) # [batch, 512]
        x = self.fc(x) # [batch, 10]

        return x

```

ResNet-18

18-layer	34-layer	50-layer	101-layer	152-layer
	7×7, 64, stride 2			
	3×3 max pool, stride 2			
$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	average pool, 1000-d fc, softmax			
1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

```

class ResNet18(nn.Module):
    def __init__(self, block, num_classes=10):
        super(ResNet18, self).__init__()

        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
        )

        self.conv2_1 = block(64, 64)
        self.conv2_2 = block(64, 64)

        self.conv3_1 = block(64, 128)
        self.conv3_2 = block(128, 128)

        self.conv4_1 = block(128, 256)
        self.conv4_2 = block(256, 256)

        self.conv5_1 = block(256, 512)
        self.conv5_2 = block(512, 512)

        self.avg_pool = nn.AdaptiveAvgPool2d((1,1)) # make output h*w = 1*1
        self.fc = nn.Linear(512, num_classes)

    def forward(self, x):
        # x: [batch, 3, 32, 32]

        x = self.conv1(x) # x: [batch, 64, 8, 8]

        x = self.conv2_1(x) # [batch, 64, 8, 8]
        x = self.conv2_2(x) # [batch, 64, 8, 8]

        x = self.conv3_1(x) # [batch, 128, 8, 8]
        x = self.conv3_2(x) # [batch, 128, 8, 8]

        x = self.conv4_1(x) # [batch, 256, 8, 8]
        x = self.conv4_2(x) # [batch, 256, 8, 8]

        x = self.conv5_1(x) # [batch, 512, 8, 8]
        x = self.conv5_2(x) # [batch, 512, 8, 8]

        x = self.avg_pool(x) # [batch, 512, 1, 1]
        x = x.view(x.size(0), -1) # [batch, 512]
        x = self.fc(x) # [batch, 10]

        return x

```

ResNet-18

18-layer	34-layer	50-layer	101-layer	152-layer
	7×7, 64, stride 2			
	3×3 max pool, stride 2			
$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
average pool, 1000-d fc, softmax				
1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

```
class ResNet18(nn.Module):
    def __init__(self, block, num_classes=10):
        super(ResNet18, self).__init__()

        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
        )

        self.conv2_1 = block(64, 64)
        self.conv2_2 = block(64, 64)

        self.conv3_1 = block(64, 128)
        self.conv3_2 = block(128, 128)

        self.conv4_1 = block(128, 256)
        self.conv4_2 = block(256, 256)

        self.conv5_1 = block(256, 512)
        self.conv5_2 = block(512, 512)

        self.avg_pool = nn.AdaptiveAvgPool2d((1,1)) # make output h*w = 1*1
        self.fc = nn.Linear(512, num_classes)
```

```
def forward(self, x):
    # x: [batch, 3, 32, 32]

    x = self.conv1(x) # x: [batch, 64, 8, 8]

    x = self.conv2_1(x) # [batch, 64, 8, 8]
    x = self.conv2_2(x) # [batch, 64, 8, 8]

    x = self.conv3_1(x) # [batch, 128, 8, 8]
    x = self.conv3_2(x) # [batch, 128, 8, 8]

    x = self.conv4_1(x) # [batch, 256, 8, 8]
    x = self.conv4_2(x) # [batch, 256, 8, 8]

    x = self.conv5_1(x) # [batch, 512, 8, 8]
    x = self.conv5_2(x) # [batch, 512, 8, 8]

    x = self.avg_pool(x) # [batch, 512, 1, 1]
    x = x.view(x.size(0), -1) # [batch, 512]
    x = self.fc(x) # [batch, 10]

    return x
```

ResNet-18

```
!pip install torchsummary
from torchsummary import summary
```

```
net = ResNet18(BasicBlock).cuda()
summary(net, batch_size=-1, input_size=(3, 32, 32), device='cuda')
```

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 64, 16, 16]	9,408
BatchNorm2d-2	[-1, 64, 16, 16]	128
ReLU-3	[-1, 64, 16, 16]	0
MaxPool2d-4	[-1, 64, 8, 8]	0
Conv2d-5	[-1, 64, 8, 8]	36,864
BatchNorm2d-6	[-1, 64, 8, 8]	128
ReLU-7	[-1, 64, 8, 8]	0
Conv2d-8	[-1, 64, 8, 8]	36,864
BatchNorm2d-9	[-1, 64, 8, 8]	128
ReLU-10	[-1, 64, 8, 8]	0
BasicBlock-11	[-1, 64, 8, 8]	0
Conv2d-12	[-1, 64, 8, 8]	36,864
BatchNorm2d-13	[-1, 64, 8, 8]	128
ReLU-14	[-1, 64, 8, 8]	0
Conv2d-15	[-1, 64, 8, 8]	36,864
BatchNorm2d-16	[-1, 64, 8, 8]	128
ReLU-17	[-1, 64, 8, 8]	0
BasicBlock-18	[-1, 64, 8, 8]	0
Conv2d-19	[-1, 128, 8, 8]	73,728
BatchNorm2d-20	[-1, 128, 8, 8]	256
ReLU-21	[-1, 128, 8, 8]	0
Conv2d-22	[-1, 128, 8, 8]	147,456

```
...
Forward/backward pass size (MB): 7.85
Params size (MB): 42.65
Estimated Total Size (MB): 50.51
-----
```

18-layer	34-layer	50-layer	101-layer
7×7, 64, stride 2			
3×3 max pool, stride 2			
$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$
$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$
$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
average pool, 1000-d fc, softmax			
1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9